SHARAT CHANDRA SAI BODDU

AI Engineer · LLMs and Applied Machine Learning

Denton, TX · (940) 629-3865 · sharatchandrasai999@gmail.com · Portfolio: sharatchandrasai999-sketch.github.io

LinkedIn: linkedin.com/in/sharat-chandra-sai-boddu-6504673a6 · GitHub: github.com/sharatchandrasai999-sketch

SUMMARY

AI engineer building practical tooling around large language models: systems that make LLM output cheaper to run and easier to trust. Built ExtractEval and TokenLens, and contributed a merged fix to simonw/llm. M.S. in AI, University of North Texas, May 2026.

TECHNICAL SKILLS

- Languages: Python, SQL, C, C++, Dart
- LLMs and GenAI: prompt engineering, LLM evaluation, agentic AI, token-cost optimization, guardrails, semantic caching, MCP
- Data and ML: PySpark, MLflow, Databricks, Delta Lake, PyTorch, scikit-learn, pandas, NumPy, Hugging Face Transformers, spaCy
- Engineering: Git, GitHub Actions, pytest, Docker, Streamlit, Linux, CI/CD, AWS (S3, Lambda, Textract)

PROJECTS

ExtractEval github.com/sharatchandrasai999-sketch/extracteval

Python, Anthropic Claude API, Streamlit, YAML, pytest, GitHub Actions

- Built ExtractEval, an open-source tool that extracts and verifies document information. New tasks need only a YAML file; models are scored on accuracy, cost, and speed.
- Ships as a command-line tool and web dashboard, runs offline without an API key, and is covered by 19 passing tests with CI on every change.

TokenLens github.com/sharatchandrasai999-sketch/tokenlens

Python, tiktoken, sentence-transformers, Streamlit, pytest, GitHub Actions

- Created TokenLens, an open-source tool that prices a prompt across 16+ AI models before it is sent. It counts tokens exactly, ranks models cheapest first, and flags calls that would go over budget.
- Ships as a CLI, Python API, and dashboard, with a cost optimizer, budget guardrails, and semantic caching; 46 tests passing.

Animal Shelter Outcome Prediction (team project)

- Team project: predicted outcomes for 174,000+ Austin Animal Center records across 8 classes, reaching 70% accuracy, 21 points above the most-common-outcome baseline.

simonw/llm (open-source contribution) github.com/simonw/llm/pull/1726

- Merged pull request #1726: fixed silent data loss when a model returned the wrong number of embeddings, so the tool raises a clear error instead. Added three regression tests; all 23 embed tests pass, with formatting and lint checks clean.

EXPERIENCE

Artificial Intelligence Intern, Thea Solutions LLC — Austin, TX (Remote) Nov 2026 – Present

- Working with AWS services (S3, Lambda, and Textract) to build AI features, using prompt engineering and GitHub Copilot to speed up development.
- Following CI/CD practices through the full development cycle, from build to deployment.

Artificial Intelligence Engineer, Zensar Technologies — Hyderabad, India Mar 2023 – Aug 2024

- Built a PySpark pipeline that cleaned raw datasets and saved them as Delta tables for efficient model training.
- Ran and compared two machine learning models in MLflow, documenting accuracy and cost to select the best model.
- Fixed a slow PySpark job with early data filtering and fewer joins, significantly reducing run time.

Information Technology Analyst, University of North Texas — Denton, TX May 2025 – May 2026

- Monitored student attendance and updated daily work schedules for student workers, keeping records accurate and coordination smooth.

Android App Development Intern, Verzeo — Hyderabad, India (Hybrid) Jan 2022 – Jun 2022

- Developed Android mobile apps, learning mobile UI/UX principles and the Android SDK, and earned the official completion certificate.

EDUCATION

University of North Texas — M.S. in Artificial Intelligence, May 2026 · GPA 3.64/4.0

DRK Institute of Science and Technology — B.Tech in Computer Science, 2020 – 2024

Certifications: LLM, Hugging Face (2026) · AI Agent, Databricks (2026) · Responsive Web Design, freeCodeCamp (2023)