

SHARAT CHANDRA SAI BODDU

AI Engineer · LLMs and Applied Machine Learning

Denton, TX · (940) 629-3865 · sharatchandrasai999@gmail.com · Portfolio: sharatchandrasai999-sketch.github.io
LinkedIn: linkedin.com/in/sharat-chandra-sai-boddu-6504673a6 · GitHub: github.com/sharatchandrasai999-sketch
Work authorization: F-1 OPT (from August 2026)

SUMMARY

M.S. in Artificial Intelligence, University of North Texas (2026). AI engineer building LLM tooling: an extraction-evaluation framework (agentic supervisor layer) and a token-cost toolkit (budget guardrails, semantic caching). Open-source contributor (simonw/llm). Seeking ML/AI engineer roles.

TECHNICAL SKILLS

- Languages: Python, SQL, C, C++
- LLMs and GenAI: Claude API, prompt engineering, RAG, LLM evaluation, agents, tiktoken, token-cost optimization, guardrails, semantic caching
- Machine learning: PyTorch, scikit-learn, pandas, NumPy, Hugging Face Transformers, spaCy
- Engineering: Git, GitHub Actions, pytest, Docker, Streamlit, REST APIs, Linux, AWS (S3)

PROJECTS

ExtractEval github.com/sharatchandrasai999-sketch/extracteval

Python, Anthropic Claude API, Streamlit, YAML, pytest, GitHub Actions

- Built ExtractEval, an open-source framework turning one YAML file (fields, rules, labels) into a complete extraction task with zero code changes for new tasks, type-aware scoring, and a multi-model leaderboard on accuracy, cost, and latency.
- Introduced an agentic supervisor layer routing extractions to auto-approve, human review, or reject by confidence; a rules engine blocks invalid results (negative totals, future dates). Auto-approved demo batch: 100% correct, zero errors escaped.
- Shipped as CLI and Streamlit dashboard with an offline provider (no API key needed); 19 pytest tests, CI green.

TokenLens github.com/sharatchandrasai999-sketch/tokenlens

Python, tiktoken, sentence-transformers, Streamlit, pytest, GitHub Actions

- Developed TokenLens, an open-source toolkit pricing prompts across 16+ models before sending, using exact tiktoken counts or an offline estimator with cheapest-first ranking to catch over-budget calls before money is spent.
- Added a three-lever optimizer (trim waste, re-route, cap output) reporting each lever's savings separately, plus a prompt linter and a budget guardrail that warns or blocks over-limit calls.
- Implemented semantic caching with a swappable embedder, scored on hit rate and false-hit rate; CLI, Python API, and dashboard; 46 tests, CI green.

Animal Shelter Outcome Prediction (team project)

- Predicted 8 shelter-outcome classes from 174,000+ Austin Animal Center records (2013–2025): 70% accuracy and 0.68 weighted F1 with a tuned k-NN, 21 points above the majority-class baseline, via engineered temporal, demographic, and medical features with SMOTE and tuned preprocessing.

OPEN SOURCE CONTRIBUTIONS

simonw/llm (open-source contribution) github.com/simonw/llm/pull/1726

- Pull request #1726: fixed silent data loss in Collection.embed_multi_with_metadata (wrong embedding counts silently dropped entries); now validates per batch and raises ValueError with expected vs actual counts.
- Wrote three regression tests (short/long batches, multi-batch boundary); 23 embed tests green, black and ruff clean.

EXPERIENCE

Student Assistant, University of North Texas — Denton, TX May 2025 – May 2026

- Coordinated scheduling, attendance, and communications for student workers.

Android App Development Intern, Verzeo — Hyderabad, India (Hybrid) Jan 2022 – Jun 2022

- Built mobile apps with Dart in a six-month industry internship; earned official certification.

EDUCATION

University of North Texas — M.S. in Artificial Intelligence, May 2026 · GPA 3.8/4.0

Coursework: ML, Deep Learning, NLP, Statistics, Big Data, Data Mining, RL

DRK Institute of Science and Technology — B.Tech in Computer Science

Certifications: Large Language Models, Hugging Face (2026) · AI Agent Accreditation, Databricks (2026)