

SHARAT CHANDRA SAI BODDU

Machine Learning Engineer | LLM / AI Engineering | Applied AI
Denton, TX • (940) 629-3865 • sharatchandrasai999@gmail.com • [LinkedIn](#) • [GitHub](#)

SUMMARY

M.S. in Artificial Intelligence candidate and applied ML / AI engineer who designs, deploys, and evaluates end-to-end machine learning and LLM systems in Python. Built open-source frameworks for LLM evaluation, AI-agent routing, and token-cost optimization plus a multi-class predictive-modeling pipeline on 174K+ records, with strengths in model development, NLP, deep learning, and production-style tooling (model deployment, CLI, dashboards, tests, CI/CD). Seeking a Machine Learning / AI Engineer role.

EDUCATION

University of North Texas — Denton, TX

Expected May 2026

M.S. in Artificial Intelligence — GPA: 3.8/4.0

- Coursework: Machine Learning, Deep Learning, NLP, Statistics & Probability, Big Data, Data Mining, Reinforcement Learning, Feature Engineering, Software Development for AI.

DRK Institute of Science and Technology — Hyderabad, India

2024

B.Tech in Computer Science

TECHNICAL SKILLS

Languages: Python, SQL, C, C++

ML / AI: PyTorch, scikit-learn, pandas, NumPy, NLP (spaCy, NLTK, Hugging Face Transformers), neural networks, deep learning, predictive modeling, imbalanced-learn (SMOTE)

LLM & GenAI: LLM APIs (Anthropic Claude), prompt engineering, fine-tuning, RAG (retrieval-augmented generation), AI agents, LLM evaluation, extraction pipelines, semantic caching, token-cost optimization, tiktoken

Tools & Deployment: Git/GitHub, Docker, Kubernetes, CI/CD (GitHub Actions), model deployment, REST APIs, AWS (S3), Streamlit, Jupyter, Linux, pytest

Data & Methods: data pipelines, big data, feature engineering, data preprocessing, Matplotlib, Seaborn

Concepts: Supervised/Unsupervised Learning, Deep Learning, Reinforcement Learning, Transformers, Model Evaluation, Class-Imbalance Handling

PROJECTS

ExtractEval — Schema-Driven LLM Extraction & Evaluation Framework

GitHub

Python, Anthropic Claude API, Streamlit, YAML, pytest, GitHub Actions

- Engineered an open-source Python framework that turns a single YAML task definition into LLM extraction, type-aware scoring, a multi-model leaderboard, and confidence-based agentic routing — adding any new extraction task with zero code changes; scored 75% exact-match on a hand-labeled set with per-field accuracy.
- Implemented a human-in-the-loop layer (auto-approve / review / reject) and a YAML validation-rules engine, driving the auto-approved batch to 100% accuracy with zero errors escaped on the demo set.

TokenLens — LLM Token Cost Calculator & Optimization Toolkit

GitHub

Python, tiktoken, Streamlit, sentence-transformers, pytest, GitHub Actions

- Developed an open-source tool estimating LLM token cost per request across 16+ models before a call, using exact tiktoken counts or an offline estimator within ~15% of exact tokenization; added a three-lever prompt optimizer (trim, model-routing, output-cap) with per-lever savings measurement.
- Built a semantic cache with a swappable embedder and a hit / false-hit threshold evaluation, turning repeated requests into zero-cost responses; shipped as a CLI, importable library, and dashboard with tests and CI.

Animal Shelter Outcome Prediction — Multi-Class ML Pipeline

Team Project

Python, scikit-learn, imbalanced-learn (SMOTE), pandas, Matplotlib, Seaborn

- Designed an end-to-end supervised ML pipeline predicting 8 shelter-outcome classes from 174K+ Austin Animal Center records (2013–2025), reaching 70% accuracy and 0.68 weighted F1 with a tuned KNN model — a 21-point accuracy gain over the majority-class baseline.
- Engineered temporal, demographic, and medical-status features and benchmarked Decision Tree, KNN, and Neural Network (MLP) models with imputation, one-hot encoding, standardization, and SMOTE to correct severe class imbalance; delivered confusion-matrix analysis surfacing welfare-critical errors.

EXPERIENCE

Student Assistant — University of North Texas

May 2025 – May 2026

Denton, TX

- Coordinated shift scheduling, attendance tracking, and staff communications (e.g., holiday and schedule notices) for student workers, ensuring accurate records and reliable coverage.

Android App Development Intern — Verzeo

Jan 2022 – Jun 2022

Hyderabad, India (Hybrid)

- Developed Android mobile applications in a 6-month industry internship using the Android SDK and Dart, applying mobile UI/UX principles and app-lifecycle management; completed with official certification.

CERTIFICATIONS

- Large Language Models — Hugging Face (2026) • AI Agent Accreditation — Databricks (2026)

ACHIEVEMENTS

- MELA Merit Scholarship (₹50,000) • Best Coordinator, Street Cause NGO (2018–19)