# SHARAT CHANDRA SAI BODDU

AI Engineer · LLMs and Applied Machine Learning

Denton, TX · (940) 629-3865 · sharatchandrasai999@gmail.com · Portfolio: sharatchandrasai999-sketch.github.io
LinkedIn: linkedin.com/in/sharat-chandra-sai-boddu-6504673a6 · GitHub: github.com/sharatchandrasai999-sketch
Work authorization: F-1 OPT (from August 2026)

## SUMMARY

I completed my M.S. in Artificial Intelligence at the University of North Texas in May 2026. I build tools for large language models. One checks whether a model pulled the right information from documents and compares models on accuracy, cost, and speed. The other checks what a prompt will cost before it is sent. I also fixed a data-loss bug in llm, a popular open-source tool.

## TECHNICAL SKILLS

- Languages: Python, SQL, C, C++
- LLMs and GenAI: Claude API, prompt engineering, LLM evaluation, agents, tiktoken, token-cost optimization, guardrails, semantic caching
- Machine learning: PyTorch, scikit-learn, pandas, NumPy, Hugging Face Transformers, spaCy
- Engineering: Git, GitHub Actions, pytest, Docker, Streamlit, REST APIs, Linux, AWS (S3)

## PROJECTS

**ExtractEval** github.com/sharatchandrasai999-sketch/extracteval

Python, Anthropic Claude API, Streamlit, YAML, pytest, GitHub Actions

- Built ExtractEval to pull key information from documents with an AI model and check whether the answers are correct. New tasks are set up in one YAML file, with no code changes. The tool scores each model on accuracy, cost, and time.
- Added a checking step for results the model is not sure about. Good results are approved on their own, doubtful results go to a person for review, and wrong results, such as negative amounts or future dates, are blocked. In the demo test, every approved result was correct.
- Available as a command-line tool and a web dashboard, and it also runs offline without an API key. 19 tests cover the main functions; all tests pass.

**TokenLens** github.com/sharatchandrasai999-sketch/tokenlens

Python, tiktoken, sentence-transformers, Streamlit, pytest, GitHub Actions

- Created TokenLens to check the cost of a prompt before it is sent to an AI model. It counts the tokens, checks the price on more than 16 models, and lists the cheapest first, so users catch a costly call before they pay.
- Added three ways to lower the cost: cut extra words, move to a cheaper model, or limit the answer length. It can also warn or stop a request that goes over budget.
- Added a store for past answers, so a repeated question does not need a new paid call. Available as a command-line tool, a Python package, and a web dashboard, with 46 tests passing.

**Animal Shelter Outcome Prediction (team project)**

- Worked in a team to predict what would happen to animals at a shelter, using 174,000+ records from the Austin Animal Center. The model sorted animals into 8 outcome groups and was correct 70% of the time, 21 points above simply picking the most common outcome.

## OPEN SOURCE CONTRIBUTIONS

**simonw/llm (open-source contribution)** github.com/simonw/llm/pull/1726

- Pull request #1726: fixed a bug where records were silently dropped when a model returned the wrong number of results. Now the tool stops and shows an error instead.
- Added three tests to make sure the bug does not come back. All 23 related tests pass.

## EXPERIENCE

**Student Assistant, University of North Texas — Denton, TX** May 2025 – May 2026

- Handled schedules, attendance, and daily communication for student workers.

**Android App Development Intern, Verzeo — Hyderabad, India (Hybrid)** Jan 2022 – Jun 2022

- Built Android apps with Dart in a six-month internship; earned a certificate.

## EDUCATION

**University of North Texas — M.S. in Artificial Intelligence, May 2026 · GPA 3.64/4.0**

Coursework: Machine Learning, Deep Learning, NLP, Big Data, Data Mining, Information Retrieval

**DRK Institute of Science and Technology — B.Tech in Computer Science**

Certifications: Large Language Models, Hugging Face (2026) · AI Agent Accreditation, Databricks (2026)